

👁 Watch 0

🍴 Fork 0

☆

Unsloth is a local UI for training and running Gemma 4, Qwen3.6, DeepSeek, Kimi, GLM and other models.

📄 Apache License 2.0

🔗 [unsloth.ai/docs](#)

📄 Code of conduct

👤 Contributing

☆ 0 stars 🍴 0 forks 👁 0 watching 🍴 1 branch 📄 0 tags ↻ Activity

🌐 Public repository · Forked from [unslothai/unsloth](#)

🍴

1 Branch

📄 0 Tags

🔍 Go to file

T

Go to file

Add file

📄

Code

...

This branch is up to date with unslothai/unsloth:main .

🍴 Contribute

🔄 Sync fork

danielhanchen and pre-commit-ci[bot]

7:48

🕒

hours ago

Bump install.sh / install.ps1 pin to unsloth>=2026.7.6 ([unslothai#7621](#))

📁 .github	CI: prove an interrupted install can neve...	13 hours ago
📁 images	images: use narrower Discord button a...	2 months ago
📁 scripts	Measure where Studio's startup time ac...	yesterday
📁 studio	Bump install.sh / install.ps1 pin to unsl...	11 hours ago
📁 tests	Studio: detect local MTP subdirectory d...	11 hours ago
📁 unsloth	Update _utils.py	11 hours ago
📁 unsloth_cli	unsloth start: harden coding-agent inst...	16 hours ago
📄 .gitattributes	Replace standalone Studio wording wit...	2 weeks ago
📄 .gitignore	Show release notes in the update popu...	yesterday
📄 .pre-commit-ci.yaml	pre-commit CI config (unslothai#3565)	8 months ago
📄 .pre-commit-config.yaml	[pre-commit.ci] pre-commit autoupdate ...	last month
📄 CHANGELOG.md	Show release notes in the update popu...	yesterday
📄 CODE_OF_CONDUCT.md	Update CODE_OF_CONDUCT.md	9 months ago
📄 CONTRIBUTING.md	docs: repository cleanup (unslothai#56...)	last month
📄 COPYING	Rename cli/ to unsloth_cli/ to fix names...	4 months ago

 LICENSE	Rename cli/ to unsloth_cli/ to fix names...	4 months ago
 MANIFEST.in	Show release notes in the update popu...	yesterday
 README.md	Studio: settings UI cleanup for Agents, ...	12 hours ago
 _changelog_build.py	Show release notes in the update popu...	yesterday
 build.sh	Show release notes in the update popu...	yesterday
 cli.py	Rename cli/ to unsloth_cli/ to fix names...	4 months ago
 install.ps1	Bump install.sh / install.ps1 pin to unsl...	11 hours ago
 install.sh	Bump install.sh / install.ps1 pin to unsl...	11 hours ago
 pyproject.toml	Update pyproject.toml	11 hours ago
 unsloth-cli.py	feat(studio): add DoRA support to studi...	last week

READMECode of conductContributingLicense



Unsloth Studio lets you run and train models locally.

[Features](#) • [News](#) • [Quickstart](#) • [Notebooks](#) • [Documentation](#)

Fine-tuning Studio

Configure and start training

ConfigureCurrent RunHistory

Model

2x Faster Training

Select base model and training method

Local Model

Hugging Face Model

Method

Hugging Face Token (Optional)

./models/my-model

unsloth/gemma-4-e4b-8

LoRA (16-bit)

hf_...

No local models found. Store path manually

Dataset

Select or upload training data

Choose dataset

Hugging Face

unsloth/radiology_mni

Subset

Train Split

default

train

Evaluation Split

None

unsloth/rad-to-logy_mni1

Hugging Face Dataset (dataset name)

Clear

Upload

View dataset

Parameters

Configure training hyperparameters

Max Steps

Use Defaults

30

Limit training to a fixed number of optimizer steps

Control Length

2048

Max sequence length for training samples

Learning Rate

8.8982

Recommended 20-40 for LoRA, 2e-5 for full fine-tune

LoRA Settings

Training Hyperparameters

Training

Monitor and control training

Training Loss

1.000

0.999

0.998

0.997

0.996

0.995

0.994

0.993

0.992

0.991

0.990

0.989

0.988

0.987

0.986

0.985

0.984

0.983

0.982

0.981

0.980

0.979

0.978

0.977

0.976

0.975

0.974

0.973

0.972

0.971

0.970

0.969

0.968

0.967

0.966

0.965

0.964

0.963

0.962

0.961

0.960

0.959

0.958

0.957

0.956

0.955

0.954

0.953

0.952

0.951

0.950

0.949

0.948

0.947

0.946

0.945

0.944

0.943

0.942

0.941

0.940

0.939

0.938

0.937

0.936

0.935

0.934

0.933

0.932

0.931

0.930

0.929

0.928

0.927

0.926

0.925

0.924

0.923

0.922

0.921

0.920

0.919

0.918

0.917

0.916

0.915

0.914

0.913

0.912

0.911

0.910

0.909

0.908

0.907

0.906

0.905

0.904

0.903

0.902

0.901

0.900

0.899

0.898

0.897

0.896

0.895

0.894

0.893

0.892

0.891

0.890

0.889

0.888

0.887

0.886

0.885

0.884

0.883

0.882

0.881

0.880

0.879

0.878

0.877

0.876

0.875

0.874

0.873

0.872

0.871

0.870

0.869

0.868

0.867

0.866

0.865

0.864

0.863

0.862

0.861

0.860

0.859

0.858

0.857

0.856

0.855

0.854

0.853

0.852

0.851

0.850

0.849

0.848

0.847

0.846

0.845

0.844

0.843

0.842

0.841

0.840

0.839

0.838

0.837

0.836

0.835

0.834

0.833

0.832

0.831

0.830

0.829

0.828

0.827

0.826

0.825

0.824

0.823

0.822

0.821

0.820

0.819

0.818

0.817

0.816

0.815

0.814

0.813

0.812

0.811

0.810

0.809

0.808

0.807

0.806

0.805

0.804

0.803

0.802

0.801

0.800

0.799

0.798

0.797

0.796

0.795

0.794

0.793

0.792

0.791

0.790

0.789

0.788

0.787

0.786

0.785

0.784

0.783

0.782

0.781

0.780

0.779

0.778

0.777

0.776

0.775

0.774

0.773

0.772

0.771

0.770

0.769

0.768

0.767

0.766

0.765

0.764

0.763

0.762

0.761

0.760

0.759

0.758

0.757

0.756

0.755

0.754

0.753

0.752

0.751

0.750

0.749

0.748

0.747

0.746

0.745

0.744

0.743

0.742

0.741

0.740

0.739

0.738

0.737

0.736

0.735

0.734

0.733

0.732

0.731

0.730

0.729

0.728

0.727

0.726

0.725

0.724

0.723

0.722

0.721

0.720

0.719

0.718

0.717

0.716

0.715

0.714

0.713

0.712

0.711

0.710

0.709

0.708

0.707

0.706

0.705

0.704

0.703

0.702

0.701

0.700

0.699

0.698

0.697

0.696

0.695

0.694

0.693

0.692

0.691

0.690

0.689

0.688

0.687

0.686

0.685

0.684

0.683

0.682

0.681

0.680

0.679

0.678

0.677

0.676

0.675

0.674

0.673

0.672

0.671

0.670

0.669

0.668

0.667

0.666

0.665

0.664

0.663

0.662

0.661

0.660

0.659

0.658

0.657

0.656

0.655

0.654

0.653

0.652

0.651

0.650

0.649

0.648

0.647

0.646

0.645

0.644

0.643

0.642

0.641

0.640

0.639

0.638

0.637

0.636

0.635

0.634

0.633

0.632

0.631

0.630

0.629

0.628

0.627

0.626

0.625

0.624

0.623

0.622

0.621

0.620

0.619

0.618

0.617

0.616

0.615

0.614

0.613

0.612

0.611

0.610

0.609

0.608

0.607

0.606

0.605

0.604

0.603

0.602

0.601

0.600

0.599

0.598

0.597

0.596

0.595

0.594

0.593

0.592

0.591

0.590

0.589

0.588

0.587

0.586

0.585

0.584

0.583

0.582

0.581

0.580

0.579

0.578

0.577

0.576

0.575

0.574

0.573

0.572

0.571

0.570

0.569

0.568

0.567

0.566

0.565

0.564

0.563

0.562

0.561

0.560

0.559

0.558

0.557

0.556

0.555

0.554

0.553

0.552

0.551

0.550

0.549

0.548

0.547

0.546

0.545

0.544

0.543

0.542

0.541

0.540

0.539

0.538

0.537

0.536

0.535

0.534

0.533

0.532

0.531

0.530

0.529

0.528

0.527

0.526

0.525

0.524

0.523

0.522

0.521

0.520

0.519

0.518

0.517

0.516

0.515

0.514

0.513

0.512

0.511

0.510

0.509

0.508

0.507

0.506

0.505

0.504

0.503

0.502

0.501

0.500

0.499

0.498

0.497

0.496

0.495

0.494

0.493

0.492

0.491

0.490

0.489

0.488

0.487

0.486

0.485

0.484

0.483

0.482

0.481

0.480

0.479

0.478

0.477

0.476

0.475

0.474

0.473

0.472

0.471

0.470

0.469

0.468

0.467

0.466

0.465

0.464

0.463

0.462

0.461

0.460

0.459

0.458

0.457

0.456

0.455

0.454

0.453

0.452

0.451

0.450

0.449

0.448

0.447

0.446

0.445

0.444

0.443

0.442

0.441

0.440

0.439

0.438

0.437

0.436

0.435

0.434

0.433

0.432

0.431

0.430

0.429

0.428

0.427

0.426

0.425

0.424

0.423

0.422

0.421

0.420

0.419

0.418

0.417

0.416

0.415

0.414

0.413

0.412

0.411

0.410

0.409

0.408

0.407

0.406

0.405

0.404

0.403

0.402

0.401

0.400

0.399

0.398

0.397

0.396

0.395

0.394

0.393

0.392

0.391

0.390

0.389

0.388

0.387

0.386

0.385

0.384

0.383

0.382

0.381

0.380

0.379

0.378

0.377

0.376

0.375

0.374

0.373

0.372

0.371

0.370

0.369

0.368

0.367

0.366

0.365

0.364

0.363

0.362

0.361

0.360

0.359

0.358

0.357

0.356

0.355

0.354

0.353

0.352

0.351

0.350

0.349

0.348

0.347

0.346

0.345

0.344

0.343

0.342

0.341

0.340

0.339

0.338

0.337

0.336

0.335

0.334

0.333

0.332

0.331

0.330

0.329

0.328

0.327

0.326

0.325

0.324

0.323

0.322

0.321

0.320

0.319

0.318

0.317

0.316

0.315

0.314

0.313

0.312

0.311

0.310

0.309

0.308

0.307

0.306

0.305

0.304

0.303

0.302

0.301

0.300

0.299

0.298

0.297

0.296

0.295

0.294

0.293

0.292

0.291

0.290

0.289

0.288

0.287

0.286

0.285

0.284

0.283

0.282

0.281

0.280

0.279

0.278

0.277

0.276

0.275

0.274

0.273

0.272

0.271

0.270

0.269

0.268

0.267

0.266

0.265

0.264

0.263

0.262

0.261

0.260

0.259

0.258

0.257

0.256

0.255

0.254

0.253

0.252

0.251

0.250

0.249

0.248

0.247

0.246

0.245

0.244

0.243

0.242

0.241

0.240

0.239

0.238

0.237

0.236

0.235

0.234

0.233

0.232

0.231

0.230

0.229

0.228

0.227

0.226

0.225

0.224

0.223

0.222

0.221

0.220

0.219

0.218

0.217

0.216

0.215

0.214

0.213

0.212

0.211

0.210

0.209

0.208

0.207

0.206

0.205

0.204

0.203

0.202

0.201

0.200

0.199

0.198

0.197

0.196

0.195

0.194

0.193

0.192

0.191

0.190

0.189

0.188

0.187

0.186

0.185

0.184

0.183

0.182

0.181

0.180

0.179

0.178

0.177

0.176

0.175

0.174

0.173

0.172

0.171

0.170

0.169

0.168

0.167

0.166

0.165

0.164

0.163

0.162

0.161

0.160

0.159

0.158

0.157

0.156

0.155

0.154

0.153

0.152

0.151

0.150

0.149

0.148

0.147

0.146

0.145

0.144

0.143

0.142

0.141

0.140

0.139

0.138

0.137

0.136

0.135

0.134

0.133

0.132

0.131

0.130

0.129

0.128

0.127

0.126

0.125

0.124

0.123

0.122

0.121

0.120

0.119

0.118

0.117

0.116

0.115

0.114

0.113

0.112

0.111

0.110

0.109

0.108

0.107

0.106

0.105

0.104

0.103

0.102

0.101

0.100

0.099

0.098

0.097

0.096

0.095

0.094

0.093

0.092

0.091

0.090

0.089

0.088

0.087

0.086

0.085

0.084

0.083

0.082

0.081

0.080

0.079

0.078

0.077

0.076

0.075

0.074

0.073

0.072

0.071

0.070

0.069

0.068

0.067

0.066

0.065

0.064

0.063

0.062

0.061

0.060

0.059

0.058

0.057

0.056

0.055

0.054

0.053

0.052

0.051

0.050

0.049

0.048

0.047

0.046

0.045

0.044

0.043

0.042

0.041

0.040

0.039

0.038

0.037

0.036

0.035

0.034

0.033

0.032

0.031

0.030

0.029

0.028

0.027

0.026

0.025

0.024

0.023

0.022

0.021

0.020

0.019

0.018

0.017

0.016

0.015

0.014

0.013

0.012

0.011

0.010

0.009

0.008

0.007

0.006

0.005

0.004

0.003

0.002

0.001

0.000

Start Training

Training Config

Upload

Save

Reset

unsloth

New Chat

Compare

Search

Train

Recipes

Export

RECENT CHATS

how to use unsloth studio?

unsloth

Run

gemma-4-268-A4B-8-GGUF GGUF - UD-Q4_KXL

2.2K / 262.1K

17

how to use unsloth studio?

Used tool: Searched "what is Unsloth Studio and how to use it"

Introducing Unsloth Stud...

Get started with Unsloth ...

GitHub - unslothai/unslo...

Unsloth - Train and Run ...

How to Run models with ...

Thought for 1 seconds

Unsloth Studio is a local, browser-based graphical user interface (GUI) designed to let you fine-tune and run Large Language Models (LLMs) on your own computer (Mac or Windows) without needing to write code.

Based on the official documentation, here is a guide on how to use its primary features:

1. Installation and Setup

Since Unsloth Studio is designed to run locally and offline:

• System Requirements: It is optimized for local hardware (Mac and Windows).

Send a message...

Think

Search

Code

LLM's can make mistakes. Double-check all responses.

Get started

macOS, Linux, WSL:

curl -fsSL https://unsloth.ai/install.sh | sh

Windows:



🔗 Community:

- [Discord](#)
- [X \(Twitter\)](#)
- [Reddit](#)

🔗 ★ Features

Unsloth Studio (Beta) lets you run and train text, [audio](#), [embedding](#), [vision](#) models on Windows, Linux and macOS.

🔗 Inference

- **Search + download + run models** including GGUF, LoRA adapters, safetensors
- **Export models:** [Save or export](#) models to GGUF, 16-bit safetensors and other formats.
- **Tool calling:** Support for [self-healing tool calling](#) and web search
- **Code execution:** lets LLMs test code in Claude artifacts and sandbox environments
- **API inference endpoint:** Deploy and run local LLMs in Claude Code, Codex tools with Unsloth
- [Auto set inference settings](#) and customize chat templates.
- We work directly with teams behind [gpt-oss](#), [Qwen3](#), [Llama 4](#), [Mistral](#), [Gemma 1-3](#), and [Phi-4](#), where we've fixed bugs that improve model accuracy.
- Chat with images, audio, PDFs, code, DOCX and more. [Connect API providers](#) (OpenAI, Anthropic) or servers (vLLM, Ollama).
- [Compare any two models](#) side by side with the same prompt.
- **OpenAI/Anthropic-compatible APIs:** Serve local models through `/v1/chat/completions`, `/v1/responses` and `/v1/messages`.
- **Connect local models to agents:** Use `unsloth start` with Claude Code, Codex, Hermes and more.
- **Web/PDF search** can read PDF papers, manuals and other PDF results.
- **GGUF hardware controls:** Choose GPUs/layers, offload MoE experts, use multi-GPU or Tensor Parallelism.
- The opt-in **MCP control endpoint** lets AI clients manage models, training, recipes and exports.

🔗 Training

- Train and RL **500+ models** up to **2x faster** with **70% less VRAM**; MoE up to **12x faster**.
- Train and run RL on [AMD GPUs](#) across Windows, WSL and Linux.
- **Data Recipes:** [Auto-create datasets](#) from **PDF, CSV, DOCX** etc. Edit data in a visual-node workflow.
- **Reinforcement Learning** uses **80% less VRAM** for GRPO, FP8 and vision RL, with 7x longer contexts.
- **Long-context training:** **3x faster**, 30% less VRAM and 500K+ context.
- Supports LoRA/QLoRA, full fine-tuning, RL, pretraining, 4-bit, 16-bit and FP8.
- Custom Triton and mathematical **kernels** built with PyTorch and Hugging Face.
- **Observability:** Monitor training live, track loss and GPU usage and customize graphs.
- [Multi-GPU](#) training is supported, with major improvements coming soon.

🔗 🚀 Unsloth Start

[Unsloth Start](#) connects [Claude Code](#), [Codex](#) and other agents to local models with one command.

Start Unsloth, load a model, open your project folder, then run:

```
unsloth start claude
```



Replace `claude` with any supported agent:

Agent	Command
Claude Code	<code>unsloth start claude</code>
OpenAI Codex	<code>unsloth start codex</code>
Hermes Agent	<code>unsloth start hermes</code>
OpenClaw	<code>unsloth start openclaw</code>
OpenCode	<code>unsloth start opencode</code>

Claude Code, Codex and OpenCode can keep their current model and use Unsloth as a local subagent:

```
unsloth start claude --as-subagent --model unsloth/model-GGUF:quant
```



 Install

Unsloth can be used in two ways: through [Unsloth Studio](#), the web UI, or through **Unsloth Core**, the code-based version. Each has different requirements.

 Unsloth Studio (web UI)

Unsloth Studio (Beta) works on **Windows, Linux, WSL** and **macOS**.

- **CPU:** Supported for Chat and Data Recipes currently
- **NVIDIA:** Training works on RTX 30/40/50, Blackwell, DGX Spark, Station and more
- **macOS:** Training, MLX and GGUF inference are ALL supported.
- **AMD:** Training, RL, chat and deployment work on Windows, WSL and Linux. [Read the AMD guide](#).
- **Vulkan:** GGUF inference is supported on [compatible GPUs, including Intel GPUs](#). Vulkan accelerates GGUF inference only; training still requires a supported PyTorch or MLX backend.
- **Multi-GPU:** Available now, with a major upgrade on the way

 macOS, Linux, WSL:

```
curl -fsSL https://unsloth.ai/install.sh | sh
```



Use the same command to update.

To force the Vulkan llama.cpp backend, set `UNSLOTH_FORCE_VULKAN=1` **before installing or updating**. The setting selects the llama.cpp binary bundle, so setting it only when launching Studio cannot replace an existing CPU bundle:

```
export UNSLOTH_FORCE_VULKAN=1
curl -fsSL https://unsloth.ai/install.sh | sh
```



🔗 Windows:

```
irm https://unsloth.ai/install.ps1 | iex
```



Use the same command to update.

To force the Vulkan llama.cpp backend, set the environment variable before running the installer or updater:

```
$env:UNSLOTH_FORCE_VULKAN=1  
irm https://unsloth.ai/install.ps1 | iex
```



Re-running the current installer replaces a previously selected CPU bundle when the backend differs. A separate Vulkan SDK is not required; the GPU driver must provide a working Vulkan runtime.

🔗 Launch

```
unsloth studio -p 8888
```



For LAN or cloud access, add `-H 0.0.0.0` (raw port only; add `--cloudflare` for a public URL). By default, Unsloth is accessible only locally.

To reach Unsloth over HTTPS, use `unsloth studio --secure`. Unsloth stays bound to localhost and is reached only through a free Cloudflare tunnel, which publishes it at a public `https://*.trycloudflare.com` URL (it fails closed if the tunnel can't start, so the raw port is never exposed). This makes Unsloth reachable from the internet, so anyone with the link and API key can use it and run code: keep your API key private (see Remote access below).

🔗 Docker

Use our [Docker image](#) `unsloth/unsloth` container. Run:

```
docker run -d -e JUPYTER_PASSWORD="mypassword" \  
  -p 8888:8888 -p 8000:8000 -p 2222:22 \  
  -v $(pwd)/work:/workspace/work \  
  --gpus all \  
  unsloth/unsloth
```



🔗 Developer, Nightly, Uninstall

To see developer, nightly and uninstallation etc. instructions, see [advanced installation](#).

🔗 Unsloth Core (code-based)

🔗 Linux, WSL:

```
curl -LsSf https://astral.sh/uv/install.sh | sh  
uv venv unsloth_env --python 3.13  
source unsloth_env/bin/activate  
uv pip install unsloth --torch-backend=auto
```



🔗 Windows:

```
winget install -e --id Python.Python.3.13
winget install --id=astral-sh.uv -e
uv venv unsloth_env --python 3.13
.\unsloth_env\Scripts\activate
uv pip install unsloth --torch-backend=auto
```



For Windows, `pip install unsloth` works only if you have PyTorch installed. Read our [Windows Guide](#). You can use the same Docker image as Unsloth Studio.

AMD, Intel:

For RTX 50x, B200, 6000 GPUs: `uv pip install unsloth --torch-backend=auto`. Read our guides for: [Blackwell](#) and [DGX Spark](#).

To install Unsloth on **AMD** and **Intel** GPUs, follow our [AMD Guide](#) and [Intel Guide](#).

Free Notebooks

Train for free with our notebooks. You can use our new [free Unsloth Studio notebook](#) to run and train models for free in a web UI. Read our [guide](#). Add dataset, run, then deploy your trained model.

Model	Free Notebooks	Performance	Memory use
Gemma 4 (E2B)	 Start for free	1.5x faster	50% less
Qwen3.5 (4B)	 Start for free	1.5x faster	60% less
gpt-oss (20B)	 Start for free	2x faster	70% less
Qwen3.5 GSPO	 Start for free	2x faster	70% less
gpt-oss (20B): GRPO	 Start for free	2x faster	80% less
Qwen3: Advanced GRPO	 Start for free	2x faster	70% less
embeddinggemma (300M)	 Start for free	2x faster	20% less
Mistral Ministral 3 (3B)	 Start for free	1.5x faster	60% less
Llama 3.1 (8B) Alpaca	 Start for free	2x faster	70% less
Llama 3.2 Conversational	 Start for free	2x faster	70% less
Orpheus-TTS (3B)	 Start for free	1.5x faster	50% less

- See all our notebooks for: [Kaggle](#), [GRPO](#), [TTS](#), [embedding](#) & [Vision](#)
- See [all our models](#) and [all our notebooks](#)
- See detailed documentation for Unsloth [here](#)

Unsloth News

- **AMD training:** Train, run RL, chat and deploy on AMD GPUs across Windows, WSL and Linux. [Guide](#)
- **GGUF hardware controls:** Choose GPU/layer placement, offload MoE experts and use multi-GPU or Tensor Parallelism. [#6414](#)
- **Local models for any agent:** Use `unsloth start` with Claude Code, Codex, Hermes, OpenCode, OpenClaw and more through Unsloth's OpenAI- and Anthropic-compatible APIs. [Guide](#)

- **MCP control endpoint:** Let compatible clients manage models, training, recipes, checkpoints and exports. [#7191](#)
- **Local inference reliability:** Resume long chats faster, recover stalled downloads and reuse existing GGUF files. [#7204](#) • [#6858](#) • [#7209](#)
- **New models:** [Qwen-AgentWorld](#), [Ornith](#), [Kimi K2.7 Code](#) and [MiniMax M3](#)
- **GLM-5.2:** Run Z.ai's 744B-parameter, 1M-context open model locally with Unsloth Dynamic GGUFs. [Guide](#)
- **DeepSeek-V4:** Run DeepSeek-V4-Flash locally with corrected multi-turn and tool-calling behavior. [Guide](#)
- **DiffusionGemma:** Run and fine-tune Google's diffusion language model with 1.8x faster inference in Unsloth Studio. [Guide](#)
- **Qwen3.6:** Run and train Qwen3.6 with MTP for 1.4-2.2x faster inference and NVFP4 quants for supported GPUs. [Guide](#)
- **Gemma 4:** Run and train Gemma 4 text, image and audio models with QAT, MTP, GGUF and MLX support. [Guide](#)
- **MCP servers:** Connect local models to files, apps, databases and external tools through Model Context Protocol. [Guide](#)
- **Connections:** Mix local models with API providers (OpenAI, Anthropic) or servers (vLLM, Ollama) in the same interface. [Guide](#)
- **Introducing Unsloth Studio:** our new web UI for running and training LLMs. [Blog](#)
- Train **MoE LLMs 12x faster** with 35% less VRAM - DeepSeek, GLM, Qwen and gpt-oss. [Blog](#)
- **Embedding models:** Unsloth now supports ~1.8-3.3x faster embedding fine-tuning. [Blog](#) • [Notebooks](#)
- New **7x longer context RL** vs. all other setups, via our new batching algorithms. [Blog](#)
- New **RoPE & MLP Triton Kernels & Padding Free + Packing:** 3x faster training & 30% less VRAM. [Blog](#)
- **500K Context:** Training a 20B model with >500K context is now possible on an 80GB GPU. [Blog](#)
- **FP8 & Vision RL:** You can now do FP8 & VLM GRPO on consumer GPUs. [FP8 Blog](#) • [Vision RL](#)

Advanced Installation

The below advanced instructions are for Unsloth Studio. For Unsloth Core advanced installation, [view our docs](#).

Developer / Nightly / Experimental installs: macOS, Linux, WSL:

The developer install builds from the `main` branch, which is the latest (nightly) source.

```
git clone https://github.com/unslothai/unsloth
cd unsloth
./install.sh --local
unsloth studio -p 8888
```



To install into an isolated location (its own virtual env, `auth/`, `studio.db`, `cache` and `llama.cpp` build), set `UNSLOTH_STUDIO_HOME` and pass it again at launch:

```
UNSLOTH_STUDIO_HOME="$PWD/.studio" ./install.sh --local
UNSLOTH_STUDIO_HOME="$PWD/.studio" unsloth studio -p 8888
```



Then to update :

```
cd unsloth && git pull
./install.sh --local
unsloth studio -p 8888
```



🔗 Developer / Nightly / Experimental installs: Windows PowerShell:

The developer install builds from the `main` branch, which is the latest (nightly) source.

```
git clone https://github.com/unslothai/unsloth.git
cd unsloth
Set-ExecutionPolicy -Scope Process -ExecutionPolicy Bypass
.\install.ps1 --local
unsloth studio -p 8888
```



To install into an isolated location (its own virtual env, `auth/`, `studio.db`, `cache` and `llama.cpp` build), set `UNSLOTH_STUDIO_HOME` and pass it again at launch:

```
$env:UNSLOTH_STUDIO_HOME="$PWD\.studio"; .\install.ps1 --local
$env:UNSLOTH_STUDIO_HOME="$PWD\.studio"; unsloth studio -p 8888
```



Then to update :

```
cd unsloth; git pull
.\install.ps1 --local
unsloth studio -p 8888
```



🔗 Remote access: --secure (HTTPS tunnel) vs raw port

By default `unsloth studio` binds to `127.0.0.1` (this machine only). To reach it from another device, pick one of:

- `--secure` (recommended): serve **only** through a free Cloudflare HTTPS link. Unsloth stays bound to localhost and the tunnel provides the public URL; it fails closed (does not start) if the tunnel can't come up, so the raw port is never exposed.

```
unsloth studio --secure -p 8888
```



- `-H 0.0.0.0` : bind the raw port on all network interfaces, reachable from anywhere on the network (subject to your firewall). It does not create a public internet URL; add `--cloudflare` to also publish an internet-reachable `https://*.trycloudflare.com` link even behind a firewall. Only use this on a network you trust.

```
unsloth studio -H 0.0.0.0 -p 8888
```



The Cloudflare tunnel is **off by default**: `-H 0.0.0.0` exposes the raw port only, not a public internet URL. Pair the wildcard bind with `--cloudflare` (`unsloth studio -H 0.0.0.0 --cloudflare`) to also publish a public `https://*.trycloudflare.com` link, or prefer `--secure` (above), which keeps the raw port private. `--cloudflare` has no effect on a loopback bind.

On a wildcard bind Unsloth works out the address to share by asking `ifconfig.me` for the public IP, then asks `check-host.net` whether that port is reachable so it can tell you if a firewall is in the way. Both contact a third party. Set `UNSLOTH_STUDIO_DISABLE_PUBLIC_CHECK=1` to skip them; the banner then shows the LAN address and no reachability line.

The first time Unsloth is published on a public URL (`--secure` or `--cloudflare`) with the auto-generated admin password still in place, it asks for a new admin password in the terminal (masked input with confirmation) before the public link goes up. Without an attached terminal it warns instead and keeps the bootstrap deadline: Unsloth shuts down after `UNSLOTH_STUDIO_BOOTSTRAP_TIMEOUT` (default 1 hour) unless the password is changed in the web UI.

For headless setups that cannot answer that prompt, set the initial admin password non-interactively with `--password` (only takes effect when no password is set yet; if one already exists it is a hard error, so rotate later with `unsloth studio reset-password`):

```
unsloth studio --secure --password 'your-strong-password'          # visible in `ps`/history
UNSLOTH_STUDIO_PASSWORD='your-strong-password' unsloth studio --secure # via env var
printf '%s\n' 'your-strong-password' | unsloth studio --secure --password - # via stdin
```

A literal `--password VALUE` is visible in the process list and shell history, so prefer the `UNSLOTH_STUDIO_PASSWORD` env var or `--password -` (stdin) for automation. This applies to any launch (public or a headless `-H 0.0.0.0 bind`), and the password is set in the parent before the server binds, so it never reaches a re-executed child process.

Server-side tools (web search, Python and terminal code execution) run as your user and are on by default. Anyone who can reach the server with the API key can run code on this machine, so keep your API key private and pass `--disable-tools` when exposing Unsloth.

🔗 Advanced launch options

Installer options can be passed as environment variables. On macOS, Linux and WSL place the variable after the pipe so the shell passes it to `sh`; on Windows set it with `$env:` before piping to `iex`.

Skip PyTorch (GGUF-only mode):

```
curl -fsSL https://unsloth.ai/install.sh | UNSLOTH_NO_TORCH=1 sh
```

```
$env:UNSLOTH_NO_TORCH=1; irm https://unsloth.ai/install.ps1 | iex
```

Skip the post-install prompt that starts Unsloth (useful for automated installs):

```
curl -fsSL https://unsloth.ai/install.sh | UNSLOTH_SKIP_AUTOSTART=1 sh
```

```
$env:UNSLOTH_SKIP_AUTOSTART=1; irm https://unsloth.ai/install.ps1 | iex
```

Pin the Python version:

```
curl -fsSL https://unsloth.ai/install.sh | UNSLOTH_PYTHON=3.12 sh
```

```
$env:UNSLOTH_PYTHON='3.12'; irm https://unsloth.ai/install.ps1 | iex
```

Install to a custom location with `UNSLOTH_STUDIO_HOME` :

```
curl -fsSL https://unsloth.ai/install.sh | UNSLOTH_STUDIO_HOME=/abs/path sh
```

```
$env:UNSLOTH_STUDIO_HOME='C:\path'; irm https://unsloth.ai/install.ps1 | iex
```



On macOS, the installer defaults to the system certificate store (`UV_SYSTEM_CERTS=1`) so uv trusts the CAs in your Keychain, needed behind TLS-inspecting proxies (Cisco Umbrella, Zscaler, etc.). Opt out with:

```
curl -fsSL https://unsloth.ai/install.sh | UV_SYSTEM_CERTS=0 sh
```



Point the frontend build at a corporate npm mirror/proxy with `UNSLOTH_NPM_REGISTRY` (for the developer install behind a firewall that blocks `registry.npmjs.org`):

```
UNSLOTH_NPM_REGISTRY=https://artifactory.example.com/api/npm/npm/ ./install.sh --local
```



```
$env:UNSLOTH_NPM_REGISTRY='https://artifactory.example.com/api/npm/npm/'; .\install.ps1 --local
```



It is threaded as `--registry` into the Unsloth frontend `npm / bun` installs; the supply-chain locks (7-day min-release-age , exact version pins) stay in force.

Cap Unsloth's native CPU thread pools on high-core hosts: `UNSLOTH_CPU_THREADS=8 unsloth studio -p 8888` .

🔗 Uninstall

The recommended way to fully remove Unsloth Studio is the matching uninstall script for your OS. It stops any running servers, removes the install dir, the launcher data dir, the desktop shortcut, and any platform-specific entries (macOS `.app bundle` + Launch Services on Mac; Start Menu, `HKCU\Software\Unsloth` registry key and user `PATH` entries on Windows):

- **MacOS, WSL, Linux:** `curl -fsSL https://raw.githubusercontent.com/unslothai/unsloth/main/scripts/uninstall.sh | sh`
- **Windows (PowerShell):** `irm https://raw.githubusercontent.com/unslothai/unsloth/main/scripts/uninstall.ps1 | iex`

If you only want to drop the install dir and keep the launcher/shortcut for a later reinstall, you can instead run `rm -rf ~/.unsloth/studio` (Mac/Linux/WSL) or `Remove-Item -Recurse -Force "$HOME\.unsloth\studio"` (Windows). The model cache at `~/.cache/huggingface` is not touched by any of these.

For more info, [see our docs](#).

🔗 Deleting model files

You can delete old model files either from the bin icon in model search or by removing the relevant cached model folder from the default Hugging Face cache directory. By default, HF uses:

- **MacOS, Linux, WSL:** `~/.cache/huggingface/hub/`
- **Windows:** `%USERPROFILE%\cache\huggingface\hub\`

🔗 ❤️ Community and Links

Type	Links
💬 Discord	Join Discord server

Type	Links
 r/unsloth Reddit	Join Reddit community
 Documentation & Wiki	Read Our Docs
 Twitter (aka X)	Follow us on X
 Our Models	Unsloth Catalog
 Blog	Read our Blogs

Citation

You can cite the Unsloth repo as follows:

```
@software{unsloth,  
  author = {Daniel Han, Michael Han and Unsloth team},  
  title = {Unsloth},  
  url = {https://github.com/unslothai/unsloth},  
  year = {2023}  
}
```



If you trained a model with 🦥 Unsloth, you can use this cool sticker!



License

Unsloth uses a dual-licensing model of Apache 2.0 and AGPL-3.0. The core Unsloth package remains licensed under [Apache 2.0](#), while certain optional components, such as the Unsloth Studio UI are licensed under the open-source license [AGPL-3.0](#).

This structure helps support ongoing Unsloth development while keeping the project open source and enabling the broader ecosystem to continue growing.

Thank You to

Releases

No releases published
[Create a new release](#)

Packages

No packages published
[Publish your first package](#)

Contributors

No contributors

Languages

Suggested workflows

Based on your tech stack



SLSA Generic generator

Generate SLSA3 provenance for your existing release workflows

By Open Source Security Foundation (OpenSSF)

Configure



Publish Python Package

Publish a Python Package to PyPI on release.

By GitHub Actions

Configure



Python package

Create and test a Python package on multiple Python versions.

By GitHub Actions

Configure

[More workflows](#)

